Enhancing Physical Consistency in Lightweight World Models

Dingrui Wang, Zhexiao Sun, et al.

The Problem: Robots Have "Fever Dreams"

Imagine you are driving a car. Before you turn the steering wheel, you quickly play a movie in your head: *"If I turn left, will I hit that truck?"* This ability to simulate the future is called a **World Model**.

Big AI models (like Sora or Genie) are great at this. They understand physics in a good manner. But they are enormous. Running them requires a massive supercomputer.

If you try to run these "dreams" on a regular robot chip (like the one in a Tesla or a small drone), you have to use a tiny, compressed model. The problem? **Tiny models get the physics wrong.** Cars in the simulation start disappearing, teleporting, or merging into each other. It is like the robot is having a fever dream—it looks trippy, but you definitely should not trust it to drive your car.

Our Goal: We want to make a tiny AI model that understands physics just as well as the giant ones, so it can run on a robot's actual onboard computer.

The Solution: Two Simple Tricks

We did not just make the model bigger. Instead, we taught it to focus. We used "human-like" intuitions to fix the physics:

1. The "Highlighter" Trick (Soft Mask)

Metaphor: Imagine reading a textbook. If you highlight the key words with a neon pen, you remember them better.

Existing methods either give the AI the raw image (which is too messy) or a strict black-and-white mask of objects (which is too rigid, like blocking out the text).

We use a **"Soft Mask."** Think of it as a glowing aura around the other cars on the road. We feed this "glow" into the AI during training. It gently tells the AI: *"Hey, pay attention to this area—there is a moving object here—but do not ignore the road lanes underneath."* This forces the small brain to focus its limited processing power on what matters: collisions and movement.

The Result: Smarter, Faster, Smaller

We tested this in a traffic simulator. Here is the punchline:

- **It is Tiny:** Our model is roughly **3x smaller** than the standard baseline.
- **It is Smarter:** Despite being smaller, it understands physics better. In our tests, the "big dumb model" often made cars vanish. Our "small smart model" kept them solid and moving correctly.
- **It Works:** We achieved a **60% boost** in overall quality scores compared to the baseline, and it runs fast enough to be used on a real laptop GPU or edge device.